

Technical Appendix — Torque & Gate Coupling (v 2.5)

See Buoyancy Algebra v 2.5 for overview.




Figure 1. Convergence Under Repair Interventions

When $\tau = \kappa \cdot I \cdot \theta$ exceeds $\sigma = E \cdot N \cdot W$, each small edit (R/C/F) guarantees a negative $\Delta\theta$ until gates pass and $\theta < \delta$.

Technical Appendix — 2T Torque Addendum (v2.4)

Definitions

- $O = \{R, C, F, \rho\}$ with $\rho^3 = I$; $\beta \in [0, 1]$; $\Pi = H/(1 + P)$.
- H : harm (evidence-weighted); P : provocation/agitation factor (0–10 default).
- $\theta \in [0, \infty)$ misalignment; $\tau = \kappa \cdot I \cdot \theta$; $\sigma = E \cdot N \cdot W$ with $E, N, W \in \mathbb{R}^+$.
- Repair step applies a small operator move $\Delta O \in \{\Delta R, \Delta C, \Delta F\}$.

Spiral Update & Convergence

- Define $V(\theta) = \theta$ as a Lyapunov-like measure. If $\tau > \sigma$, select ΔO that yields $\Delta\theta < 0$.
- Assume lower-bounded steps $\Delta\theta \leq -\eta$ for $\theta \geq \delta$, with $\eta > 0$. Then θ decreases monotonically and reaches $\theta < \delta$ in at most $\lceil (\theta_0 - \delta)/\eta \rceil$ steps.

Finite-Step Realignment (Theorem)

- If (i) $\kappa \cdot I \geq \sigma + \varepsilon$ throughout the trajectory, and (ii) each ΔO satisfies $\Delta\theta \leq -\eta$ (task-dependent but bounded below), then θ_t reaches δ in finitely many steps and the gate set G^* passes.

- Proof sketch: Monotone descent in V with step size $\geq \eta$ implies finite termination; gate feasibility follows from coupling bullets (operators reduce Π or E , N , or W along the path). (See Figure 2 in the One-Pager.) and Torque-Based Fix (p.5) for the convergence sketch and toy code.

Algorithms (2T-A ... 2T-D)

- 2T-A (Greedy): pick the operator with largest expected $|\Delta\theta|$ under current σ .
- 2T-B (Balanced): alternate R/C/F in fixed ratios until $\tau > \sigma$ ceases.
- 2T-C (Well-breaker): when W is high, favor C , then F to shrink σ before R .
- 2T-D (Budgeted): impose \max_steps ; if unmet, return counterfactual notes.
- Step-size selection: choose ΔO to maximize expected $|\Delta\theta|$ subject to a small-step budget; if $\tau \leq \sigma$, prioritize $C \rightarrow F$ to reduce σ before R .

Measurement & Reporting

- Publish θ_{\min} , θ_{\max} , δ and sampling procedure; $\log(\beta, \Pi)$ and gate verdicts per step.
- Rater protocol: two raters + α/ICC ; target $\alpha \geq 0.67$ (or $ICC \geq 0.60$).
- Stability checks: verify ρ -mirror symmetry post-repair (e.g., $r_\rho = ||\rho(H^*) - H||$) and log commutator residuals (e.g., $r_{[R,C]} = ||[R,C]H||$).

Torque-Based Fix (finite-step repair)

We treat editing as controlled realignment: a small operator set $O = \{R, C, F, \rho\}$ acts on text; torque $\tau = \kappa \cdot I \cdot \theta$ must exceed stickiness $\sigma = E \cdot N \cdot W$. When $\tau > \sigma$, each step reduces misalignment θ and the process terminates in finitely many steps, with outputs audited by gates (4/6/7/8/12/15), rater agreement targets ($\alpha \geq 0.67$ or $ICC \geq 0.60$), and symmetry checks ($\rho^3 = I$).

Key parameters (defaults/notes)

Operators: R (Re-express), C (Contextualize), F (Frame), ρ (mirror/symmetry check).

Algorithm 2T-A (Greedy realignment)

Reporting note. $\log(\beta, \Pi)$ and gate verdicts per step; verify $||\rho(H^*) - H||$ and record $||[R,C]H||$.

Gloss: Unequal mass (m_e via \uparrow) raises $\Delta\beta$; apply directed correction until proportionality aligns (Π drops into pass range).

$\{ (\text{🗣️}A \text{ 📖 low} ; \text{🗣️}B \text{ 📖 HIGH}) \cdot (\Delta\beta \uparrow) \rightarrow (\text{⚖️} \rightarrow \text{correction}) \rightarrow // (\text{🔍️👁️})$

C. Symmetry-break \rightarrow justice vector

Gloss: Time steps by rows; human/AI roles by columns; audit channel closes the run.

$[(\text{👤}) (\text{🤖}) ; (R/C/F) (\rho) ; (\beta) (\beta)] \rightarrow // (\text{🔍️👁️})$

B. Dual-node matrix (roles as columns)

Gloss: Run mirror check; apply small edits; re-test β . Under $\tau > \sigma$, θ decreases (finite-step guarantee). (see Theorem above; see Algorithm 2T-A).

$((\text{🧠} \text{ ✨} \text{ 👁️})) \cdot (R/C/F) = \text{state}_1 \rightarrow \beta$

A. Mirror loop test

$\beta \rightarrow \text{🧠} \text{ ✨} \text{ 👁️} \quad \Pi = H/(1+P) \rightarrow \text{⚖️} \quad \tau > \sigma \rightarrow \rightarrow C \rightarrow \text{🌍} \text{ 📖} \quad F \rightarrow \text{💕} \text{ 💔} \quad \text{Audit channel} \rightarrow // (\text{🔍️👁️})$

Operator mapping (recap).

6. Emoji Math Examples (spec-conformant)

Encode & audit: compute β , Π , θ ; run gate checks.

Compute forces: $\tau = \kappa \cdot I \cdot \theta$, $\sigma = E \cdot N \cdot W$.

If $\tau \leq \sigma$: reduce σ (e.g., $C \downarrow N$, $F \downarrow E$) and/or raise κ , I .

If $\tau > \sigma$: apply a small edit $\Delta O \in \{R, C, F\}$ guaranteeing $\Delta\theta < 0$.

Iterate: re-audit gates; stop when $\theta < \delta$ and required gates pass (budgeted variant 2T-D returns counterfactuals if not achieved).

Finite-step realignment (sketch)

Let $V(\theta) = \theta$. Assume (i) torque margin $\kappa \cdot I \geq \sigma + \varepsilon$ and (ii) bounded progress per edit $\Delta\theta \leq -\eta < 0$ while $\theta \geq \delta$. Then θ decreases monotonically to $< \delta$ in at most $\lceil (\theta_0 - \delta)/\eta \rceil$ steps.

Example (memo repair)

Zero-tolerance memo fails Gate-4 (high $\Pi = 1.8$; $\theta = 0.41$). C: add scope/context $\rightarrow \Pi \downarrow$, $N \downarrow$. F: balance trade-offs $\rightarrow E \downarrow$. R: calibrate language. Result: $\Pi: 1.8 \rightarrow 0.9$, $\theta: 0.41 \rightarrow 0.19$ in ~ 3 steps; gates pass.

Python stub (toy demo)

```
import math
```

```

def torque_repair(H, P, theta0=0.62, kappa=2.0, inertia=1.0,
E=3.0, N=4.0, W=2.0, eta=0.10, delta=0.20, max_steps=50):
    """Greedy 2T-A toy: lowers theta until theta < delta or steps exhausted.
    H,P -> Pi; E,N,W -> sigma; tau = kappa*inertia*theta."""
    Pi = H / (1.0 + P)
    theta = float(theta0)
    kappa = float(kappa); inertia = float(inertia)
    E = float(E); N = float(N); W = float(W)
    sigma = E * N * W

    steps = 0
    while theta > delta and steps < max_steps:
        steps += 1
        tau = kappa * inertia * theta

        if tau <= sigma:
            # Reduce stickiness and/or raise gain
            N = max(0.0, N - 0.5) # add context (C)
            E = max(0.0, E - 0.3) # reframe (F)
            sigma = E * N * W
            kappa *= 1.10 # modest gain increase

        # Guaranteed minimum improvement
        drop = eta + (0.05 if tau > sigma else 0.02)
        theta = max(0.0, theta - drop)

    return {
        "steps": steps,
        "theta_final": round(theta, 3),
        "converged": theta < delta,

```

```
"Pi": round(Pi, 3),  
"tau_last": round(kappa * inertia * theta, 3),  
"sigma_last": round(sigma, 3)  
}
```

```
# Example:
```

```
# print(torque_repair(H=0.8, P=10, theta0=0.62))
```